

Convolutional neural networks

Deep learning course for industry

Mitko Veta


Eindhoven University of Technology
Department of Biomedical Engineering

2020

Learning goals

- ▶ Demonstrate how deep neural networks can be modified to be more suitable for image data.
- ▶ Introduce the Keras neural networks API

Images as inputs to a neural network



5x5 image




3 hidden
neurons




1 output
neurons

Images as inputs to a neural network




5x5 image



3 hidden
neurons


1 output
neurons

Images as inputs to a neural network




The biases $w_{i,0}$ that are not shown.


Images as inputs to a neural network



The number of parameters explodes with larger image sizes



The number of parameters explodes with larger image sizes




$$\# \text{ parameters} = (\text{height} \times \text{width} \times \# \text{ channels} + 1) \times \# \text{ neurons}$$

The "+1" comes from the biases $w_{i,0}$.

Towards convolutional neural networks


Example (1-D image for simplicity): 5×1 input image, 3 hidden neurons.




full connectivity: 15
parameters

Towards convolutional neural networks

Example (1-D image for simplicity): 5×1 input image, 3 hidden neurons.




full connectivity: 15
parameters




sparse connectivity:
9 parameters

Towards convolutional neural networks


Example (1-D image for simplicity): 5×1 input image, 3 hidden neurons.



full connectivity: 15
parameters



sparse connectivity:
9 parameters




shared weights: 3
parameters

Note: the poor biases are again, ignored, but there are three of them in each case

Towards convolutional neural networks

Let the outputs of the three neurons be $\sigma(a_1), \sigma(a_2), \sigma(a_3)$. Then:




$$a_1 = x_1 w_1 + x_2 w_2 + x_3 w_3$$

$$a_2 = x_2 w_1 + x_3 w_2 + x_4 w_3$$

$$a_3 = x_3 w_1 + x_4 w_2 + x_5 w_3$$

Towards convolutional neural networks

Let the outputs of the three neurons be $\sigma(a_1), \sigma(a_2), \sigma(a_3)$. Then:



$$a_1 = x_1 w_1 + x_2 w_2 + x_3 w_3$$


$$a_2 = x_2 w_1 + x_3 w_2 + x_4 w_3$$

$$a_3 = x_3 w_1 + x_4 w_2 + x_5 w_3$$


$$[a_1, a_2, a_3] = [x_1, x_2, x_3, x_4, x_5] * [w_3, w_2, w_1]$$

, where $*$ is the convolution operator, thus a **convolutional layer**.

Motivation (or rather a justification)




sparse connectivity
motivation: the features
appear locally

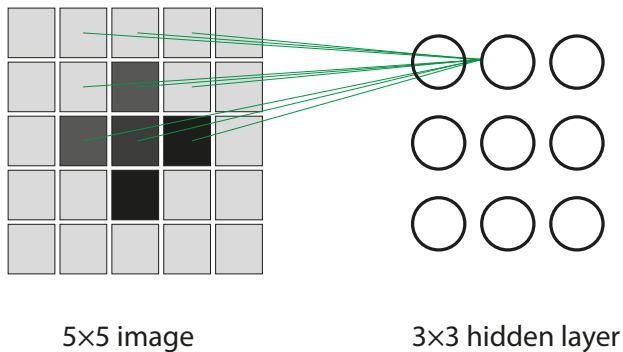


shared weights
motivation: the features
repeat throughout the image

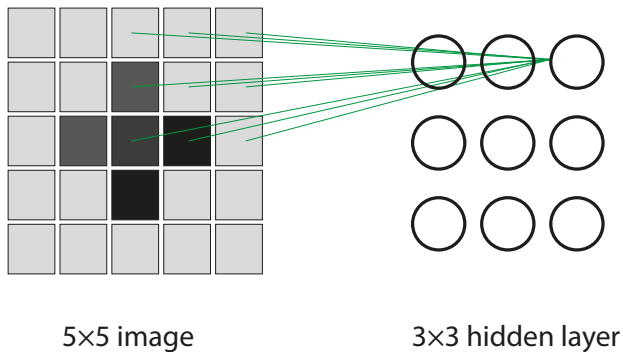
Towards convolutional neural networks in 2D



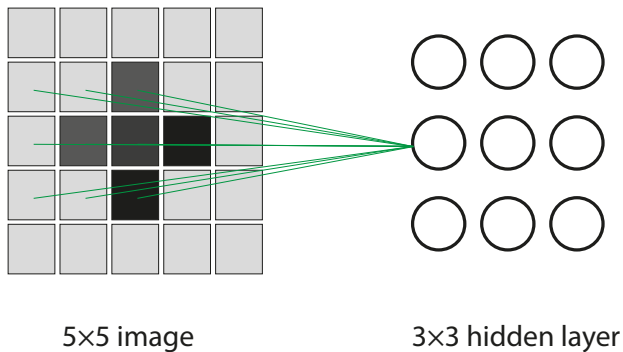
Towards convolutional neural networks in 2D




Towards convolutional neural networks in 2D




Towards convolutional neural networks in 2D




Towards convolutional neural networks in 2D




Towards convolutional neural networks in 2D



Adding a second feature map




5x5 image




3x3x2 hidden layer
(2 feature maps)

Adding a second feature map



Adding a second feature map



Convolution with padding

Figure source: https://github.com/vdumoulin/conv_arithmetic


Computing the output size

$$\text{output size} = \frac{\text{input size} - \text{kernel size} + 2 \times \text{padding}}{\text{stride}} + 1$$

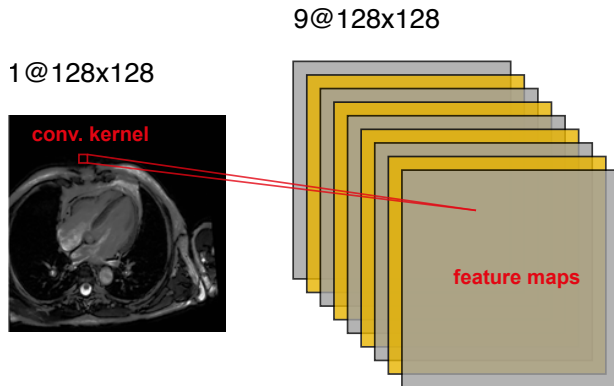
In this example: input size = 5, kernel size = 3, padding = 1, stride = 1.
The output size is $(5 - 3 + 2 \times 1)/1 + 1 = 5$.

Motivation (or rather justification) for CNNs

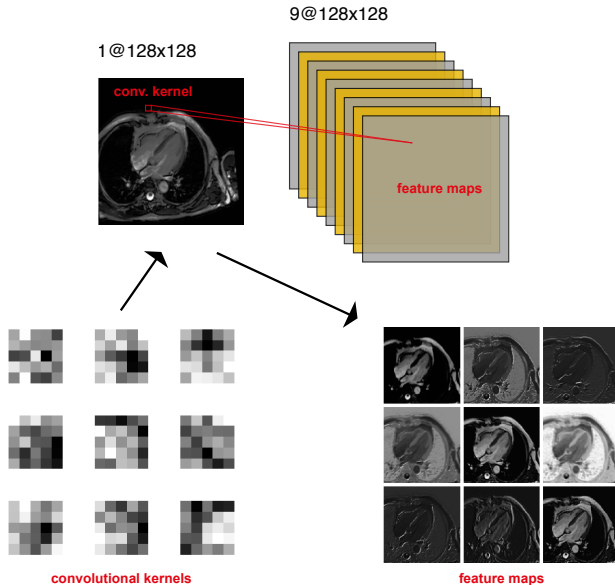
The features of interest can appear at different locations in the image.



Kernels and feature maps



Kernels and feature maps



Motivation (or rather a consequence) for deep CNNs

The network learns low-level features in the first layers, and builds up towards more complex features in the deeper layers: intensity → edges and colour blobs → junctions → shapes → etc.




Figure source: nvidia.com

Equivariance and invariance to translation

The convolutional layers are **equivariant with translation**: as the input is translated, the output is translated in a predictable manner.

Equivariance and invariance to translation

The convolutional layers are **equivariant with translation**: as the input is translated, the output is translated in a predictable manner.

A desired property of neural networks for classification is **invariance**: as the input is translated, the output remains the same.

Partial translational invariance of CNNs is achieved with the max-pooling operator.


Note: there are other types of invariance e.g. rotational.

Max-pooling


A max-pool with a 2×2 kernel stride and size 2 (most common form) will reduce the image size by 2 in each dimension (a useful side-effect).

A “typical” CNN architecture for 2D image classification

Note that the convolution is a linear operation so non-linearities (such as ReLU) are still needed.




Training CNNs



Example by Andrej Karpathy.

What is Keras?

Keras is a high-level neural networks API, written in Python and capable of running on top of TensorFlow, CNTK, or Theano.



What is Keras?

Two API implementations: keras.io (reference implementation) and from Tensorflow.

```
import keras
import tensorflow as tf
```

```
layer = keras.layers.Dense # used in slides
layer = tf.keras.layers.Dense # used in the exercises
```

Implementing neural networks in Keras

The core data structure of Keras is a **model**, a way to organize layers.

The simplest type of model is the **Sequential** model, a linear stack of layers. Covers majority of use cases.


For more complex architectures, you should use the Keras **functional API**, which allows to build arbitrary graphs of layers. Covers almost all use cases.

Functional API

Most common use cases:

- ▶ Models with multiple input and/or multiple outputs
- ▶ Models with shared layers


Example implementation



placeholder for RGB images


```
inputs = keras.Input(shape=[64, 64, 1])
```

Example implementation




```
conv1 = keras.layers.Conv2D(  
    filters      = 16,  
    kernel_size  = 3,  
    strides      = (1, 1),  
    padding      = "same",  
    activation   = "relu")
```


Example implementation




```
maxPool1 = keras.layers.MaxPool2D(  
    pool_size = (2, 2),  
    strides   = (2, 2),  
    padding   = "valid")
```

Example implementation




```
conv2 = keras.layers.Conv2D(  
    filters      = 32,  
    kernel_size  = 3,  
    strides      = (1, 1),  
    padding      = "same",  
    activation   = "relu")
```

Example implementation




```
maxPool2 = keras.layers.MaxPool2D(  
    pool_size = (2, 2),  
    strides    = (2, 2),  
    padding    = "valid")
```

Example implementation




```
dense = keras.layers.Dense(  
    units      = 100,  
    activation = "ReLU")
```

Example implementation



```
output = keras.layers.Dense(  
    units      = 1,  
    activation = "sigmoid")
```

Example implementation



Making the connection:

```
x = maxPool1(conv1(inputs))  
x = maxPool2(conv2(x))  
x = flatten(x)  
x = dense(x)  
outputs = output(x)
```

Training a model

```
model = keras.Model(inputs = inputs , outputs = outputs)

loss = keras.losses.categorical_crossentropy
optimizer = keras.optimizers.SGD(lr=0.01, momentum=0.9)

model.compile(loss=loss , optimizer=optimizer)
```

Training a model

```
model = keras.Model(inputs = inputs , outputs = outputs)

loss = keras.losses.categorical_crossentropy
optimizer = keras.optimizers.SGD(lr=0.01, momentum=0.9)

model.compile(loss=loss , optimizer=optimizer)

# x_train and y_train are Numpy arrays
model.fit(x_train , y_train , epochs=5, batch_size=32)
```


Training a model

Or, feed batches manually in a training loop:

```
for i in range(0, num_iterations):  
    x_batch, y_batch = some_batch_generator(i)  
    model.train_on_batch(x_batch, y_batch)
```

Model evaluation

```
loss_and_metrics = model.evaluate(x_test, y_test,  
                                  batch_size=128)
```

```
classes = model.predict(x_test, batch_size=128)
```

Summary

- ▶ Compared to fully connected neural networks, convolutional neural networks have sparse connectivity and weight sharing, which makes them suitable for image data.
- ▶ The Keras neural networks API allows for user-friendly implementation and training of (convolutional) neural networks.